

# Coding Companion for Intuitive Deep Learning Part 1 (Annotated)

Code Author: Joseph Lee Wei En

Additional Annotations in Italics (or additional Python comments) by Kayla Bandy

Date Annotated: 11/25/22

In this notebook, we'll go through the code for the coding companion for Intuitive Deep Learning Part 1 ([Part 1a](#), [Part 1b](#)) to create your very first neural network to predict whether the house price is below or above median value. We will go through the following in this notebook:

- Exploring and Processing the Data
- Building and Training our Neural Network
- Visualizing Loss and Accuracy
- Adding Regularization to our Neural Network

The code is annotated throughout the notebook and you simply need to download the dataset [here](#), put the dataset in the same folder as this notebook and run the code cells below. Note that the results you get might differ slightly from the blogpost as there is a degree of randomness in the way we split our dataset as well as the initialization of our neural network.

*This Jupyter Notebook will take us through 3 neural network examples, with the first one being relatively simple. The 2nd example shows exaggerated overfitting while the 3rd example takes the 2nd example and implements some common strategies to reduce over-fitting.*

## Exploring and Processing the Data

We first have to read in the CSV file that we've been given. We'll use a package called pandas for that:

```
In [1]: import pandas as pd
```

```
In [2]: df = pd.read_csv('housepricedata.csv')
```

```
In [3]: #View some of the dataset to see the structure of the data
#Measurements are in sq ft
df
```

Out[3]:	LotArea	OverallQual	OverallCond	TotalBsmtSF	FullBath	HalfBath	BedroomAbvGr	TotRmsAbvGrd	Firep
0	8450	7	5	856	2	1	3	8	
1	9600	6	8	1262	2	0	3	6	
2	11250	7	5	920	2	1	3	6	
3	9550	7	5	756	1	0	3	7	
4	14260	8	5	1145	2	1	4	9	
...	...	...	...	...	...	...	...	...	...
1455	7917	6	5	953	2	1	3	7	
1456	13175	6	6	1542	2	0	3	7	
1457	9042	7	9	1152	2	0	4	9	
1458	9717	5	6	1078	1	0	2	5	
1459	9937	5	6	1256	1	1	3	6	

1460 rows × 11 columns

The dataset that we have now is in what we call a pandas dataframe. To convert it to an array, simply access its values:

```
In [4]: #Create an array from the DataFrame values
dataset = df.values
```

```
In [5]: #View some of the dataset to see the structure of the data
dataset
```

```
Out[5]: array([[ 8450,    7,    5, ...,    0,   548,    1],
 [ 9600,    6,    8, ...,    1,   460,    1],
 [11250,    7,    5, ...,    1,   608,    1],
 ...,
 [ 9042,    7,    9, ...,    2,   252,    1],
 [ 9717,    5,    6, ...,    0,   240,    0],
 [ 9937,    5,    6, ...,    0,   276,    0]], dtype=int64)
```

Now, we split the dataset into our input features and the label we wish to predict.

```
In [6]: #Select the first 10 columns into input features (X)
#Select the 11th column into the predictor label (Y)
X = dataset[:,0:10]
Y = dataset[:,10]
```

Normalizing our data is very important, as we want the input features to be on the same order of magnitude to make our training easier. We'll use a min-max scaler from scikit-learn which scales our data to be between 0 and 1.

*Rescaling variables is conceptually similar to Z-scores from statistics, where a value is displayed in relation to the mean of a group of values.*

```
In [7]: from sklearn import preprocessing
```

```
In [8]: #This function will scale or convert the dataset's features  
#to values between 0 and 1  
min_max_scaler = preprocessing.MinMaxScaler()  
X_scale = min_max_scaler.fit_transform(X)
```

```
In [9]: #View how the data changed after scaling  
X_scale
```

```
Out[9]: array([[0.0334198 , 0.66666667, 0.5          , ..., 0.5          , 0.          ,  
                0.3864598 ],  
                [0.03879502, 0.55555556, 0.875          , ..., 0.33333333, 0.33333333,  
                0.32440056],  
                [0.04650728, 0.66666667, 0.5          , ..., 0.33333333, 0.33333333,  
                0.42877292],  
                ...,  
                [0.03618687, 0.66666667, 1.          , ..., 0.58333333, 0.66666667,  
                0.17771509],  
                [0.03934189, 0.44444444, 0.625          , ..., 0.25          , 0.          ,  
                0.16925247],  
                [0.04037019, 0.44444444, 0.625          , ..., 0.33333333, 0.          ,  
                0.19464034]])
```

Lastly, we wish to set aside some parts of our dataset for a validation set and a test set. We use the function `train_test_split` from `scikit-learn` to do that.

*We use the `train_test_split` function twice since we want to sub-split the validation and test set.*

```
In [10]: from sklearn.model_selection import train_test_split
```

```
In [11]: #Split total data into 70% training and 30% validation & test  
X_train, X_val_and_test, Y_train, Y_val_and_test = train_test_split(X_scale, Y, test_size=0.3)
```

```
In [12]: #Split validation & test dataset into 50% validation and test each  
#of the original 30% dataset (15% each)  
X_val, X_test, Y_val, Y_test = train_test_split(X_val_and_test, Y_val_and_test, test_size=0.5)
```

```
In [13]: #View the dimensions of the resulting datasets  
print(X_train.shape, X_val.shape, X_test.shape, Y_train.shape, Y_val.shape, Y_test.shape)  
  
(1022, 10) (219, 10) (219, 10) (1022,) (219,) (219,)
```

## Building and Training Our First Neural Network

We will be using Keras to build our architecture. Let's import the code from Keras that we will need to use:

```
In [14]: from keras.models import Sequential  
         from keras.layers import Dense
```

We will be using the `Sequential` model, which means that we merely need to describe the layers above in sequence. Our neural network has three layers:

- Hidden layer 1: 32 neurons, ReLU activation
- Hidden layer 2: 32 neurons, ReLU activation
- Output Layer: 1 neuron, Sigmoid activation



Epoch 1/100  
32/32 [=====] - 1s 10ms/step - loss: 0.7114 - acc: 0.5176 - val\_loss: 0.7186 - val\_acc: 0.4612

Epoch 2/100  
32/32 [=====] - 0s 3ms/step - loss: 0.7011 - acc: 0.5176 - val\_loss: 0.7063 - val\_acc: 0.4612

Epoch 3/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6947 - acc: 0.5176 - val\_loss: 0.6982 - val\_acc: 0.4612

Epoch 4/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6900 - acc: 0.5166 - val\_loss: 0.6918 - val\_acc: 0.4612

Epoch 5/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6862 - acc: 0.5157 - val\_loss: 0.6864 - val\_acc: 0.4521

Epoch 6/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6827 - acc: 0.5372 - val\_loss: 0.6819 - val\_acc: 0.5205

Epoch 7/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6796 - acc: 0.6037 - val\_loss: 0.6776 - val\_acc: 0.6484

Epoch 8/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6764 - acc: 0.6556 - val\_loss: 0.6735 - val\_acc: 0.6986

Epoch 9/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6733 - acc: 0.6898 - val\_loss: 0.6692 - val\_acc: 0.7352

Epoch 10/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6699 - acc: 0.7094 - val\_loss: 0.6643 - val\_acc: 0.7717

Epoch 11/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6660 - acc: 0.7329 - val\_loss: 0.6594 - val\_acc: 0.7808

Epoch 12/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6622 - acc: 0.7436 - val\_loss: 0.6544 - val\_acc: 0.7808

Epoch 13/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6584 - acc: 0.7485 - val\_loss: 0.6497 - val\_acc: 0.7854

Epoch 14/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6544 - acc: 0.7593 - val\_loss: 0.6450 - val\_acc: 0.7991

Epoch 15/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6504 - acc: 0.7603 - val\_loss: 0.6403 - val\_acc: 0.7991

Epoch 16/100  
32/32 [=====] - 0s 4ms/step - loss: 0.6463 - acc: 0.7593 - val\_loss: 0.6352 - val\_acc: 0.8082

Epoch 17/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6421 - acc: 0.7701 - val\_loss: 0.6304 - val\_acc: 0.7900

Epoch 18/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6376 - acc: 0.7652 - val\_loss: 0.6253 - val\_acc: 0.7945

Epoch 19/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6331 - acc: 0.7691 - val\_loss: 0.6196 - val\_acc: 0.8174

Epoch 20/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6284 - acc: 0.7740 - val\_loss: 0.6140 - val\_acc: 0.8219

Epoch 21/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6233 - acc: 0.7750 - val\_loss: 0.

6078 - val\_acc: 0.8265  
Epoch 22/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6182 - acc: 0.7818 - val\_loss: 0.  
6017 - val\_acc: 0.8311  
Epoch 23/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6127 - acc: 0.7847 - val\_loss: 0.  
5957 - val\_acc: 0.8311  
Epoch 24/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6071 - acc: 0.7828 - val\_loss: 0.  
5893 - val\_acc: 0.8265  
Epoch 25/100  
32/32 [=====] - 0s 3ms/step - loss: 0.6012 - acc: 0.7886 - val\_loss: 0.  
5828 - val\_acc: 0.8174  
Epoch 26/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5949 - acc: 0.7867 - val\_loss: 0.  
5748 - val\_acc: 0.8402  
Epoch 27/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5883 - acc: 0.7896 - val\_loss: 0.  
5672 - val\_acc: 0.8447  
Epoch 28/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5816 - acc: 0.7896 - val\_loss: 0.  
5593 - val\_acc: 0.8447  
Epoch 29/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5746 - acc: 0.7886 - val\_loss: 0.  
5511 - val\_acc: 0.8447  
Epoch 30/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5676 - acc: 0.7945 - val\_loss: 0.  
5441 - val\_acc: 0.8402  
Epoch 31/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5604 - acc: 0.7945 - val\_loss: 0.  
5344 - val\_acc: 0.8447  
Epoch 32/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5528 - acc: 0.7945 - val\_loss: 0.  
5257 - val\_acc: 0.8493  
Epoch 33/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5451 - acc: 0.7955 - val\_loss: 0.  
5173 - val\_acc: 0.8539  
Epoch 34/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5371 - acc: 0.7945 - val\_loss: 0.  
5071 - val\_acc: 0.8584  
Epoch 35/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5294 - acc: 0.8043 - val\_loss: 0.  
4992 - val\_acc: 0.8584  
Epoch 36/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5213 - acc: 0.8004 - val\_loss: 0.  
4900 - val\_acc: 0.8584  
Epoch 37/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5130 - acc: 0.8092 - val\_loss: 0.  
4820 - val\_acc: 0.8493  
Epoch 38/100  
32/32 [=====] - 0s 3ms/step - loss: 0.5050 - acc: 0.8014 - val\_loss: 0.  
4716 - val\_acc: 0.8676  
Epoch 39/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4972 - acc: 0.8131 - val\_loss: 0.  
4631 - val\_acc: 0.8676  
Epoch 40/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4891 - acc: 0.8229 - val\_loss: 0.  
4544 - val\_acc: 0.8630  
Epoch 41/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4811 - acc: 0.8141 - val\_loss: 0.  
4442 - val\_acc: 0.8767  
Epoch 42/100

32/32 [=====] - 0s 3ms/step - loss: 0.4734 - acc: 0.8229 - val\_loss: 0.  
4362 - val\_acc: 0.8767  
Epoch 43/100  
32/32 [=====] - 0s 2ms/step - loss: 0.4660 - acc: 0.8288 - val\_loss: 0.  
4285 - val\_acc: 0.8767  
Epoch 44/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4584 - acc: 0.8317 - val\_loss: 0.  
4206 - val\_acc: 0.8767  
Epoch 45/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4512 - acc: 0.8327 - val\_loss: 0.  
4120 - val\_acc: 0.8767  
Epoch 46/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4437 - acc: 0.8356 - val\_loss: 0.  
4031 - val\_acc: 0.8767  
Epoch 47/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4371 - acc: 0.8376 - val\_loss: 0.  
3947 - val\_acc: 0.8904  
Epoch 48/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4301 - acc: 0.8493 - val\_loss: 0.  
3920 - val\_acc: 0.8721  
Epoch 49/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4241 - acc: 0.8376 - val\_loss: 0.  
3828 - val\_acc: 0.8904  
Epoch 50/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4180 - acc: 0.8434 - val\_loss: 0.  
3760 - val\_acc: 0.8904  
Epoch 51/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4123 - acc: 0.8464 - val\_loss: 0.  
3699 - val\_acc: 0.8904  
Epoch 52/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4064 - acc: 0.8493 - val\_loss: 0.  
3649 - val\_acc: 0.8904  
Epoch 53/100  
32/32 [=====] - 0s 3ms/step - loss: 0.4013 - acc: 0.8454 - val\_loss: 0.  
3589 - val\_acc: 0.8950  
Epoch 54/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3961 - acc: 0.8434 - val\_loss: 0.  
3524 - val\_acc: 0.8950  
Epoch 55/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3916 - acc: 0.8503 - val\_loss: 0.  
3502 - val\_acc: 0.8904  
Epoch 56/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3868 - acc: 0.8513 - val\_loss: 0.  
3467 - val\_acc: 0.8858  
Epoch 57/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3824 - acc: 0.8513 - val\_loss: 0.  
3408 - val\_acc: 0.8904  
Epoch 58/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3779 - acc: 0.8532 - val\_loss: 0.  
3384 - val\_acc: 0.8813  
Epoch 59/100  
32/32 [=====] - 0s 4ms/step - loss: 0.3744 - acc: 0.8542 - val\_loss: 0.  
3287 - val\_acc: 0.8950  
Epoch 60/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3707 - acc: 0.8571 - val\_loss: 0.  
3283 - val\_acc: 0.8904  
Epoch 61/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3672 - acc: 0.8581 - val\_loss: 0.  
3248 - val\_acc: 0.8858  
Epoch 62/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3641 - acc: 0.8542 - val\_loss: 0.  
3209 - val\_acc: 0.8904

Epoch 63/100  
32/32 [=====] - 0s 2ms/step - loss: 0.3605 - acc: 0.8571 - val\_loss: 0.  
3157 - val\_acc: 0.8995  
Epoch 64/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3572 - acc: 0.8591 - val\_loss: 0.  
3130 - val\_acc: 0.8950  
Epoch 65/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3546 - acc: 0.8620 - val\_loss: 0.  
3101 - val\_acc: 0.8904  
Epoch 66/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3518 - acc: 0.8611 - val\_loss: 0.  
3089 - val\_acc: 0.8858  
Epoch 67/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3488 - acc: 0.8630 - val\_loss: 0.  
3055 - val\_acc: 0.8858  
Epoch 68/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3460 - acc: 0.8620 - val\_loss: 0.  
3064 - val\_acc: 0.8858  
Epoch 69/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3440 - acc: 0.8669 - val\_loss: 0.  
2996 - val\_acc: 0.8858  
Epoch 70/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3415 - acc: 0.8659 - val\_loss: 0.  
2982 - val\_acc: 0.8858  
Epoch 71/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3398 - acc: 0.8679 - val\_loss: 0.  
2958 - val\_acc: 0.8858  
Epoch 72/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3371 - acc: 0.8679 - val\_loss: 0.  
2942 - val\_acc: 0.8858  
Epoch 73/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3354 - acc: 0.8650 - val\_loss: 0.  
2928 - val\_acc: 0.8858  
Epoch 74/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3328 - acc: 0.8650 - val\_loss: 0.  
2873 - val\_acc: 0.8904  
Epoch 75/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3312 - acc: 0.8708 - val\_loss: 0.  
2889 - val\_acc: 0.8904  
Epoch 76/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3294 - acc: 0.8669 - val\_loss: 0.  
2847 - val\_acc: 0.8858  
Epoch 77/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3276 - acc: 0.8689 - val\_loss: 0.  
2847 - val\_acc: 0.8858  
Epoch 78/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3256 - acc: 0.8708 - val\_loss: 0.  
2853 - val\_acc: 0.8858  
Epoch 79/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3248 - acc: 0.8708 - val\_loss: 0.  
2822 - val\_acc: 0.8904  
Epoch 80/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3228 - acc: 0.8718 - val\_loss: 0.  
2811 - val\_acc: 0.8904  
Epoch 81/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3208 - acc: 0.8699 - val\_loss: 0.  
2778 - val\_acc: 0.8904  
Epoch 82/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3194 - acc: 0.8708 - val\_loss: 0.  
2755 - val\_acc: 0.8858  
Epoch 83/100  
32/32 [=====] - 0s 3ms/step - loss: 0.3175 - acc: 0.8728 - val\_loss: 0.



```

2729 - val_acc: 0.8904
Epoch 84/100
32/32 [=====] - 0s 3ms/step - loss: 0.3174 - acc: 0.8708 - val_loss: 0.
2747 - val_acc: 0.8904
Epoch 85/100
32/32 [=====] - 0s 3ms/step - loss: 0.3156 - acc: 0.8689 - val_loss: 0.
2729 - val_acc: 0.8904
Epoch 86/100
32/32 [=====] - 0s 3ms/step - loss: 0.3143 - acc: 0.8748 - val_loss: 0.
2728 - val_acc: 0.8904
Epoch 87/100
32/32 [=====] - 0s 3ms/step - loss: 0.3133 - acc: 0.8718 - val_loss: 0.
2735 - val_acc: 0.8904
Epoch 88/100
32/32 [=====] - 0s 3ms/step - loss: 0.3115 - acc: 0.8718 - val_loss: 0.
2747 - val_acc: 0.8950
Epoch 89/100
32/32 [=====] - 0s 3ms/step - loss: 0.3104 - acc: 0.8738 - val_loss: 0.
2689 - val_acc: 0.8950
Epoch 90/100
32/32 [=====] - 0s 3ms/step - loss: 0.3092 - acc: 0.8718 - val_loss: 0.
2675 - val_acc: 0.8950
Epoch 91/100
32/32 [=====] - 0s 3ms/step - loss: 0.3078 - acc: 0.8708 - val_loss: 0.
2711 - val_acc: 0.8950
Epoch 92/100
32/32 [=====] - 0s 3ms/step - loss: 0.3068 - acc: 0.8689 - val_loss: 0.
2669 - val_acc: 0.8950
Epoch 93/100
32/32 [=====] - 0s 3ms/step - loss: 0.3056 - acc: 0.8767 - val_loss: 0.
2622 - val_acc: 0.8904
Epoch 94/100
32/32 [=====] - 0s 3ms/step - loss: 0.3043 - acc: 0.8708 - val_loss: 0.
2652 - val_acc: 0.8950
Epoch 95/100
32/32 [=====] - 0s 5ms/step - loss: 0.3023 - acc: 0.8748 - val_loss: 0.
2692 - val_acc: 0.8995
Epoch 96/100
32/32 [=====] - 0s 3ms/step - loss: 0.3030 - acc: 0.8748 - val_loss: 0.
2618 - val_acc: 0.8995
Epoch 97/100
32/32 [=====] - 0s 3ms/step - loss: 0.3015 - acc: 0.8757 - val_loss: 0.
2669 - val_acc: 0.8950
Epoch 98/100
32/32 [=====] - 0s 3ms/step - loss: 0.3005 - acc: 0.8767 - val_loss: 0.
2598 - val_acc: 0.8995
Epoch 99/100
32/32 [=====] - 0s 3ms/step - loss: 0.2998 - acc: 0.8708 - val_loss: 0.
2602 - val_acc: 0.8995
Epoch 100/100
32/32 [=====] - 0s 3ms/step - loss: 0.2976 - acc: 0.8816 - val_loss: 0.
2566 - val_acc: 0.8904

```

Evaluating our data on the test set:

```
In [18]: #Get the model accuracy
model.evaluate(X_test, Y_test)[1]
```

```
7/7 [=====] - 0s 2ms/step - loss: 0.2890 - acc: 0.8767
```

```
Out[18]: 0.8767123222351074
```


# Visualizing Loss and Accuracy

Import the relevant package we need to do the visualization

```
In [19]: import matplotlib.pyplot as plt
```


We want to visualize the training loss and the validation loss like this:

```
In [20]: plt.plot(hist.history['loss'])
plt.plot(hist.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.show()
```



We can also visualize the training accuracy and the validation accuracy like this:

```
In [21]: plt.plot(hist.history['acc'])
plt.plot(hist.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='lower right')
plt.show()
```



## Adding Regularization to our Neural Network

We'll train a model which will overfit, which we call Model 2. This might take a few minutes.

*There are 2 additional hidden layers in this model, all layers have 1000 neurons, and Adam is a common optimizer that reaches the lower loss function faster.*

```
In [22]: model_2 = Sequential([
    Dense(1000, activation='relu', input_shape=(10,)),
    Dense(1000, activation='relu'),
    Dense(1000, activation='relu'),
    Dense(1000, activation='relu'),
    Dense(1, activation='sigmoid'),
])
model_2.compile(optimizer='adam',
                loss='binary_crossentropy',
                metrics=['acc'])
hist_2 = model_2.fit(X_train, Y_train,
                    batch_size=32, epochs=100,
                    validation_data=(X_val, Y_val))
```

Epoch 1/100  
32/32 [=====] - 1s 26ms/step - loss: 0.5294 - acc: 0.7476 - val\_loss: 0.3373 - val\_acc: 0.9041  
Epoch 2/100  
32/32 [=====] - 1s 20ms/step - loss: 0.3616 - acc: 0.8532 - val\_loss: 0.3039 - val\_acc: 0.9087  
Epoch 3/100  
32/32 [=====] - 1s 20ms/step - loss: 0.3144 - acc: 0.8601 - val\_loss: 0.3093 - val\_acc: 0.8858  
Epoch 4/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2811 - acc: 0.8816 - val\_loss: 0.2830 - val\_acc: 0.9041  
Epoch 5/100  
32/32 [=====] - 1s 20ms/step - loss: 0.3284 - acc: 0.8415 - val\_loss: 0.2947 - val\_acc: 0.9087  
Epoch 6/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2550 - acc: 0.8865 - val\_loss: 0.3111 - val\_acc: 0.9132  
Epoch 7/100  
32/32 [=====] - 1s 21ms/step - loss: 0.2432 - acc: 0.8865 - val\_loss: 0.3269 - val\_acc: 0.8995  
Epoch 8/100  
32/32 [=====] - 1s 23ms/step - loss: 0.2297 - acc: 0.9022 - val\_loss: 0.2939 - val\_acc: 0.9132  
Epoch 9/100  
32/32 [=====] - 1s 21ms/step - loss: 0.2350 - acc: 0.8992 - val\_loss: 0.3749 - val\_acc: 0.8904  
Epoch 10/100  
32/32 [=====] - 1s 21ms/step - loss: 0.2535 - acc: 0.8914 - val\_loss: 0.2871 - val\_acc: 0.9041  
Epoch 11/100  
32/32 [=====] - 1s 21ms/step - loss: 0.2114 - acc: 0.9090 - val\_loss: 0.4917 - val\_acc: 0.8721  
Epoch 12/100  
32/32 [=====] - 1s 21ms/step - loss: 0.2333 - acc: 0.8992 - val\_loss: 0.3536 - val\_acc: 0.8995  
Epoch 13/100  
32/32 [=====] - 1s 24ms/step - loss: 0.2314 - acc: 0.9022 - val\_loss: 0.2979 - val\_acc: 0.9132  
Epoch 14/100  
32/32 [=====] - 1s 22ms/step - loss: 0.2190 - acc: 0.8973 - val\_loss: 0.2991 - val\_acc: 0.9087  
Epoch 15/100  
32/32 [=====] - 1s 21ms/step - loss: 0.2083 - acc: 0.8992 - val\_loss: 0.3205 - val\_acc: 0.9041  
Epoch 16/100  
32/32 [=====] - 1s 21ms/step - loss: 0.2155 - acc: 0.9080 - val\_loss: 0.3263 - val\_acc: 0.9224  
Epoch 17/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2136 - acc: 0.9110 - val\_loss: 0.3027 - val\_acc: 0.8950  
Epoch 18/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2384 - acc: 0.9100 - val\_loss: 0.3553 - val\_acc: 0.8995  
Epoch 19/100  
32/32 [=====] - 1s 19ms/step - loss: 0.2167 - acc: 0.9080 - val\_loss: 0.3415 - val\_acc: 0.9041  
Epoch 20/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2009 - acc: 0.9080 - val\_loss: 0.3329 - val\_acc: 0.9087  
Epoch 21/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1990 - acc: 0.9159 - val\_loss:

0.3145 - val\_acc: 0.9087  
Epoch 22/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2094 - acc: 0.9070 - val\_loss:  
0.3045 - val\_acc: 0.9041  
Epoch 23/100  
32/32 [=====] - 1s 21ms/step - loss: 0.2149 - acc: 0.9041 - val\_loss:  
0.2879 - val\_acc: 0.9041  
Epoch 24/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2093 - acc: 0.9129 - val\_loss:  
0.3155 - val\_acc: 0.9041  
Epoch 25/100  
32/32 [=====] - 1s 19ms/step - loss: 0.2134 - acc: 0.9061 - val\_loss:  
0.3559 - val\_acc: 0.8995  
Epoch 26/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2538 - acc: 0.8865 - val\_loss:  
0.3146 - val\_acc: 0.8995  
Epoch 27/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2192 - acc: 0.9100 - val\_loss:  
0.3336 - val\_acc: 0.9041  
Epoch 28/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2006 - acc: 0.9119 - val\_loss:  
0.3198 - val\_acc: 0.9087  
Epoch 29/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1807 - acc: 0.9178 - val\_loss:  
0.3383 - val\_acc: 0.9041  
Epoch 30/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1907 - acc: 0.9207 - val\_loss:  
0.3349 - val\_acc: 0.9041  
Epoch 31/100  
32/32 [=====] - 1s 22ms/step - loss: 0.1819 - acc: 0.9188 - val\_loss:  
0.3411 - val\_acc: 0.9087  
Epoch 32/100  
32/32 [=====] - 1s 25ms/step - loss: 0.1745 - acc: 0.9266 - val\_loss:  
0.3748 - val\_acc: 0.9087  
Epoch 33/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1822 - acc: 0.9198 - val\_loss:  
0.3658 - val\_acc: 0.9041  
Epoch 34/100  
32/32 [=====] - 1s 24ms/step - loss: 0.1796 - acc: 0.9256 - val\_loss:  
0.3255 - val\_acc: 0.9132  
Epoch 35/100  
32/32 [=====] - 1s 24ms/step - loss: 0.2068 - acc: 0.9159 - val\_loss:  
0.3620 - val\_acc: 0.9087  
Epoch 36/100  
32/32 [=====] - 1s 24ms/step - loss: 0.1926 - acc: 0.9227 - val\_loss:  
0.3941 - val\_acc: 0.9087  
Epoch 37/100  
32/32 [=====] - 1s 23ms/step - loss: 0.1767 - acc: 0.9237 - val\_loss:  
0.3963 - val\_acc: 0.8904  
Epoch 38/100  
32/32 [=====] - 1s 22ms/step - loss: 0.2040 - acc: 0.9100 - val\_loss:  
0.2992 - val\_acc: 0.8904  
Epoch 39/100  
32/32 [=====] - 1s 21ms/step - loss: 0.2122 - acc: 0.9090 - val\_loss:  
0.3165 - val\_acc: 0.8950  
Epoch 40/100  
32/32 [=====] - 1s 22ms/step - loss: 0.1826 - acc: 0.9286 - val\_loss:  
0.4209 - val\_acc: 0.8995  
Epoch 41/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1729 - acc: 0.9256 - val\_loss:  
0.3817 - val\_acc: 0.9224  
Epoch 42/100

32/32 [=====] - 1s 21ms/step - loss: 0.1744 - acc: 0.9305 - val\_loss: 0.3944 - val\_acc: 0.9178  
Epoch 43/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1997 - acc: 0.9129 - val\_loss: 0.4712 - val\_acc: 0.8904  
Epoch 44/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1752 - acc: 0.9266 - val\_loss: 0.4340 - val\_acc: 0.8995  
Epoch 45/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1799 - acc: 0.9295 - val\_loss: 0.3568 - val\_acc: 0.9087  
Epoch 46/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1757 - acc: 0.9178 - val\_loss: 0.3704 - val\_acc: 0.9041  
Epoch 47/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2018 - acc: 0.9080 - val\_loss: 0.3627 - val\_acc: 0.9087  
Epoch 48/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1656 - acc: 0.9354 - val\_loss: 0.4203 - val\_acc: 0.9087  
Epoch 49/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1715 - acc: 0.9198 - val\_loss: 0.3763 - val\_acc: 0.9041  
Epoch 50/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1525 - acc: 0.9315 - val\_loss: 0.4688 - val\_acc: 0.9087  
Epoch 51/100  
32/32 [=====] - 1s 22ms/step - loss: 0.1810 - acc: 0.9188 - val\_loss: 0.4263 - val\_acc: 0.8904  
Epoch 52/100  
32/32 [=====] - 1s 20ms/step - loss: 0.2162 - acc: 0.9041 - val\_loss: 0.3500 - val\_acc: 0.8950  
Epoch 53/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1852 - acc: 0.9237 - val\_loss: 0.3489 - val\_acc: 0.8995  
Epoch 54/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1643 - acc: 0.9325 - val\_loss: 0.4327 - val\_acc: 0.9041  
Epoch 55/100  
32/32 [=====] - 1s 22ms/step - loss: 0.1675 - acc: 0.9256 - val\_loss: 0.4887 - val\_acc: 0.8995  
Epoch 56/100  
32/32 [=====] - 1s 22ms/step - loss: 0.1590 - acc: 0.9384 - val\_loss: 0.4765 - val\_acc: 0.9087  
Epoch 57/100  
32/32 [=====] - 1s 22ms/step - loss: 0.1687 - acc: 0.9266 - val\_loss: 0.4386 - val\_acc: 0.9041  
Epoch 58/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1707 - acc: 0.9237 - val\_loss: 0.5437 - val\_acc: 0.8858  
Epoch 59/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1682 - acc: 0.9247 - val\_loss: 0.4130 - val\_acc: 0.9178  
Epoch 60/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1592 - acc: 0.9374 - val\_loss: 0.4360 - val\_acc: 0.8858  
Epoch 61/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1598 - acc: 0.9335 - val\_loss: 0.4758 - val\_acc: 0.9087  
Epoch 62/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1536 - acc: 0.9325 - val\_loss: 0.5729 - val\_acc: 0.9132

Epoch 63/100  
32/32 [=====] - 1s 19ms/step - loss: 0.1513 - acc: 0.9335 - val\_loss: 0.4354 - val\_acc: 0.8904  
Epoch 64/100  
32/32 [=====] - 1s 19ms/step - loss: 0.1492 - acc: 0.9374 - val\_loss: 0.5650 - val\_acc: 0.8995  
Epoch 65/100  
32/32 [=====] - 1s 19ms/step - loss: 0.1549 - acc: 0.9325 - val\_loss: 0.4152 - val\_acc: 0.9132  
Epoch 66/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1694 - acc: 0.9217 - val\_loss: 0.3672 - val\_acc: 0.8995  
Epoch 67/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1561 - acc: 0.9344 - val\_loss: 0.5215 - val\_acc: 0.9132  
Epoch 68/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1578 - acc: 0.9295 - val\_loss: 0.5394 - val\_acc: 0.8950  
Epoch 69/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1519 - acc: 0.9413 - val\_loss: 0.3956 - val\_acc: 0.9087  
Epoch 70/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1415 - acc: 0.9423 - val\_loss: 0.5476 - val\_acc: 0.8995  
Epoch 71/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1534 - acc: 0.9335 - val\_loss: 0.4615 - val\_acc: 0.8904  
Epoch 72/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1397 - acc: 0.9393 - val\_loss: 0.5581 - val\_acc: 0.9132  
Epoch 73/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1693 - acc: 0.9344 - val\_loss: 0.4443 - val\_acc: 0.8995  
Epoch 74/100  
32/32 [=====] - 1s 21ms/step - loss: 0.1677 - acc: 0.9217 - val\_loss: 0.3535 - val\_acc: 0.9087  
Epoch 75/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1525 - acc: 0.9374 - val\_loss: 0.4147 - val\_acc: 0.9041  
Epoch 76/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1301 - acc: 0.9403 - val\_loss: 0.4984 - val\_acc: 0.9132  
Epoch 77/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1304 - acc: 0.9472 - val\_loss: 0.4668 - val\_acc: 0.9178  
Epoch 78/100  
32/32 [=====] - 1s 20ms/step - loss: 0.1548 - acc: 0.9354 - val\_loss: 0.5246 - val\_acc: 0.9087  
Epoch 79/100  
32/32 [=====] - 1s 22ms/step - loss: 0.1431 - acc: 0.9413 - val\_loss: 0.4890 - val\_acc: 0.9041  
Epoch 80/100  
32/32 [=====] - 1s 23ms/step - loss: 0.1361 - acc: 0.9462 - val\_loss: 0.5592 - val\_acc: 0.9087  
Epoch 81/100  
32/32 [=====] - 1s 22ms/step - loss: 0.1288 - acc: 0.9403 - val\_loss: 0.6431 - val\_acc: 0.8995  
Epoch 82/100  
32/32 [=====] - 1s 23ms/step - loss: 0.1429 - acc: 0.9374 - val\_loss: 0.6051 - val\_acc: 0.9087  
Epoch 83/100  
32/32 [=====] - 1s 22ms/step - loss: 0.1137 - acc: 0.9521 - val\_loss:

```

0.5751 - val_acc: 0.8995
Epoch 84/100
32/32 [=====] - 1s 21ms/step - loss: 0.1193 - acc: 0.9540 - val_loss:
0.6031 - val_acc: 0.9132
Epoch 85/100
32/32 [=====] - 1s 22ms/step - loss: 0.1378 - acc: 0.9403 - val_loss:
0.5880 - val_acc: 0.9087
Epoch 86/100
32/32 [=====] - 1s 22ms/step - loss: 0.1141 - acc: 0.9521 - val_loss:
0.5663 - val_acc: 0.8995
Epoch 87/100
32/32 [=====] - 1s 21ms/step - loss: 0.1278 - acc: 0.9521 - val_loss:
0.7668 - val_acc: 0.8950
Epoch 88/100
32/32 [=====] - 1s 22ms/step - loss: 0.1901 - acc: 0.9237 - val_loss:
0.4090 - val_acc: 0.9041
Epoch 89/100
32/32 [=====] - 1s 22ms/step - loss: 0.1449 - acc: 0.9472 - val_loss:
0.5495 - val_acc: 0.8995
Epoch 90/100
32/32 [=====] - 1s 21ms/step - loss: 0.1178 - acc: 0.9599 - val_loss:
0.5126 - val_acc: 0.9041
Epoch 91/100
32/32 [=====] - 1s 20ms/step - loss: 0.1144 - acc: 0.9550 - val_loss:
0.4947 - val_acc: 0.9132
Epoch 92/100
32/32 [=====] - 1s 20ms/step - loss: 0.1066 - acc: 0.9589 - val_loss:
0.7052 - val_acc: 0.8950
Epoch 93/100
32/32 [=====] - 1s 20ms/step - loss: 0.1136 - acc: 0.9569 - val_loss:
0.5853 - val_acc: 0.9269
Epoch 94/100
32/32 [=====] - 1s 21ms/step - loss: 0.1130 - acc: 0.9569 - val_loss:
0.5815 - val_acc: 0.8995
Epoch 95/100
32/32 [=====] - 1s 21ms/step - loss: 0.1332 - acc: 0.9501 - val_loss:
0.4411 - val_acc: 0.9041
Epoch 96/100
32/32 [=====] - 1s 21ms/step - loss: 0.1340 - acc: 0.9452 - val_loss:
0.4967 - val_acc: 0.9087
Epoch 97/100
32/32 [=====] - 1s 20ms/step - loss: 0.1480 - acc: 0.9393 - val_loss:
0.5045 - val_acc: 0.9087
Epoch 98/100
32/32 [=====] - 1s 22ms/step - loss: 0.1082 - acc: 0.9569 - val_loss:
0.5534 - val_acc: 0.9132
Epoch 99/100
32/32 [=====] - 1s 20ms/step - loss: 0.1006 - acc: 0.9667 - val_loss:
0.4671 - val_acc: 0.9041
Epoch 100/100
32/32 [=====] - 1s 21ms/step - loss: 0.0911 - acc: 0.9648 - val_loss:
0.6131 - val_acc: 0.9087

```

Let's do the same visualization to see what overfitting looks like in terms of the loss and accuracy.

*The loss for the Validation data continues to increase while the loss for the Training data reduces, which is a sign of overfitting.*


```

In [23]: plt.plot(hist_2.history['loss'])
plt.plot(hist_2.history['val_loss'])
plt.title('Model loss')


```



```
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.show()
```



```
In [24]: plt.plot(hist_2.history['acc'])
plt.plot(hist_2.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='lower right')
plt.show()
```



To address the overfitting we see in Model 2, we'll incorporate L2 regularization and dropout in our third model here (Model 3).

```
In [25]: from keras.layers import Dropout
         from keras import regularizers
```

```
In [26]: #Change the model code to include Dropout and regularization to
         #include the the squared values of the parameters into the Loss function
         #Dropout of 0.3 means a %30 probability of dropping out during training
model_3 = Sequential([
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01), input_shape=(10,)),
    Dropout(0.3),
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    Dropout(0.3),
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    Dropout(0.3),
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    Dropout(0.3),
    Dense(1, activation='sigmoid', kernel_regularizer=regularizers.l2(0.01)),
])
```

```
In [27]: model_3.compile(optimizer='adam',
                        loss='binary_crossentropy',
                        metrics=['acc'])
hist_3 = model_3.fit(X_train, Y_train,
                    batch_size=32, epochs=100,
                    validation_data=(X_val, Y_val))
```

Epoch 1/100  
32/32 [=====] - 2s 43ms/step - loss: 14.2309 - acc: 0.6311 - val\_loss: 3.8815 - val\_acc: 0.8265  
Epoch 2/100  
32/32 [=====] - 1s 35ms/step - loss: 1.6598 - acc: 0.8395 - val\_loss: 0.5977 - val\_acc: 0.9087  
Epoch 3/100  
32/32 [=====] - 1s 35ms/step - loss: 0.5772 - acc: 0.8483 - val\_loss: 0.4699 - val\_acc: 0.9041  
Epoch 4/100  
32/32 [=====] - 1s 35ms/step - loss: 0.5342 - acc: 0.8513 - val\_loss: 0.4638 - val\_acc: 0.8995  
Epoch 5/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4841 - acc: 0.8757 - val\_loss: 0.4374 - val\_acc: 0.8950  
Epoch 6/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4845 - acc: 0.8708 - val\_loss: 0.4274 - val\_acc: 0.9087  
Epoch 7/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4741 - acc: 0.8777 - val\_loss: 0.4397 - val\_acc: 0.8904  
Epoch 8/100  
32/32 [=====] - 1s 33ms/step - loss: 0.4857 - acc: 0.8620 - val\_loss: 0.4230 - val\_acc: 0.8995  
Epoch 9/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4877 - acc: 0.8611 - val\_loss: 0.4284 - val\_acc: 0.8950  
Epoch 10/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4836 - acc: 0.8679 - val\_loss: 0.4210 - val\_acc: 0.9041  
Epoch 11/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4576 - acc: 0.8777 - val\_loss: 0.4129 - val\_acc: 0.9087  
Epoch 12/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4896 - acc: 0.8581 - val\_loss: 0.4341 - val\_acc: 0.8904  
Epoch 13/100  
32/32 [=====] - 1s 33ms/step - loss: 0.4683 - acc: 0.8630 - val\_loss: 0.4096 - val\_acc: 0.9041  
Epoch 14/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4566 - acc: 0.8708 - val\_loss: 0.4140 - val\_acc: 0.8950  
Epoch 15/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4566 - acc: 0.8708 - val\_loss: 0.4188 - val\_acc: 0.9041  
Epoch 16/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4476 - acc: 0.8767 - val\_loss: 0.4058 - val\_acc: 0.9087  
Epoch 17/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4398 - acc: 0.8757 - val\_loss: 0.4249 - val\_acc: 0.8904  
Epoch 18/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4564 - acc: 0.8796 - val\_loss: 0.4115 - val\_acc: 0.8858  
Epoch 19/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4543 - acc: 0.8669 - val\_loss: 0.4569 - val\_acc: 0.8858  
Epoch 20/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4495 - acc: 0.8806 - val\_loss: 0.4610 - val\_acc: 0.8813  
Epoch 21/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4690 - acc: 0.8601 - val\_loss:

0.4646 - val\_acc: 0.8813  
Epoch 22/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4707 - acc: 0.8630 - val\_loss: 0.4409 - val\_acc: 0.9087  
Epoch 23/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4510 - acc: 0.8708 - val\_loss: 0.4104 - val\_acc: 0.9087  
Epoch 24/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4537 - acc: 0.8767 - val\_loss: 0.4560 - val\_acc: 0.8813  
Epoch 25/100  
32/32 [=====] - 1s 38ms/step - loss: 0.4375 - acc: 0.8845 - val\_loss: 0.4281 - val\_acc: 0.9087  
Epoch 26/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4829 - acc: 0.8620 - val\_loss: 0.4389 - val\_acc: 0.8858  
Epoch 27/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4499 - acc: 0.8796 - val\_loss: 0.4016 - val\_acc: 0.9041  
Epoch 28/100  
32/32 [=====] - 1s 31ms/step - loss: 0.4399 - acc: 0.8718 - val\_loss: 0.4014 - val\_acc: 0.9087  
Epoch 29/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4529 - acc: 0.8767 - val\_loss: 0.4197 - val\_acc: 0.8995  
Epoch 30/100  
32/32 [=====] - 1s 38ms/step - loss: 0.4439 - acc: 0.8679 - val\_loss: 0.4011 - val\_acc: 0.9087  
Epoch 31/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4423 - acc: 0.8816 - val\_loss: 0.4140 - val\_acc: 0.8858  
Epoch 32/100  
32/32 [=====] - 1s 38ms/step - loss: 0.4429 - acc: 0.8787 - val\_loss: 0.4132 - val\_acc: 0.8995  
Epoch 33/100  
32/32 [=====] - 1s 38ms/step - loss: 0.4447 - acc: 0.8708 - val\_loss: 0.3995 - val\_acc: 0.9087  
Epoch 34/100  
32/32 [=====] - 1s 40ms/step - loss: 0.4389 - acc: 0.8748 - val\_loss: 0.4197 - val\_acc: 0.8858  
Epoch 35/100  
32/32 [=====] - 1s 38ms/step - loss: 0.4375 - acc: 0.8816 - val\_loss: 0.4396 - val\_acc: 0.8904  
Epoch 36/100  
32/32 [=====] - 1s 38ms/step - loss: 0.4434 - acc: 0.8708 - val\_loss: 0.4157 - val\_acc: 0.8904  
Epoch 37/100  
32/32 [=====] - 1s 40ms/step - loss: 0.4669 - acc: 0.8679 - val\_loss: 0.4062 - val\_acc: 0.8950  
Epoch 38/100  
32/32 [=====] - 1s 38ms/step - loss: 0.4455 - acc: 0.8699 - val\_loss: 0.4051 - val\_acc: 0.9132  
Epoch 39/100  
32/32 [=====] - 1s 38ms/step - loss: 0.4375 - acc: 0.8767 - val\_loss: 0.4302 - val\_acc: 0.8995  
Epoch 40/100  
32/32 [=====] - 1s 39ms/step - loss: 0.4491 - acc: 0.8826 - val\_loss: 0.4018 - val\_acc: 0.8995  
Epoch 41/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4494 - acc: 0.8767 - val\_loss: 0.4135 - val\_acc: 0.8995  
Epoch 42/100

32/32 [=====] - 1s 38ms/step - loss: 0.4337 - acc: 0.8787 - val\_loss: 0.4457 - val\_acc: 0.8813  
Epoch 43/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4519 - acc: 0.8640 - val\_loss: 0.4093 - val\_acc: 0.9087  
Epoch 44/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4700 - acc: 0.8699 - val\_loss: 0.4321 - val\_acc: 0.8858  
Epoch 45/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4425 - acc: 0.8708 - val\_loss: 0.4108 - val\_acc: 0.8904  
Epoch 46/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4415 - acc: 0.8806 - val\_loss: 0.4005 - val\_acc: 0.9041  
Epoch 47/100  
32/32 [=====] - 1s 40ms/step - loss: 0.4441 - acc: 0.8787 - val\_loss: 0.4071 - val\_acc: 0.9132  
Epoch 48/100  
32/32 [=====] - 1s 37ms/step - loss: 0.4451 - acc: 0.8826 - val\_loss: 0.4176 - val\_acc: 0.8858  
Epoch 49/100  
32/32 [=====] - 1s 37ms/step - loss: 0.4470 - acc: 0.8767 - val\_loss: 0.4132 - val\_acc: 0.8858  
Epoch 50/100  
32/32 [=====] - 1s 41ms/step - loss: 0.4515 - acc: 0.8718 - val\_loss: 0.4007 - val\_acc: 0.9132  
Epoch 51/100  
32/32 [=====] - 1s 37ms/step - loss: 0.4356 - acc: 0.8777 - val\_loss: 0.4028 - val\_acc: 0.9041  
Epoch 52/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4369 - acc: 0.8826 - val\_loss: 0.4029 - val\_acc: 0.8858  
Epoch 53/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4539 - acc: 0.8728 - val\_loss: 0.4120 - val\_acc: 0.8858  
Epoch 54/100  
32/32 [=====] - 1s 38ms/step - loss: 0.4421 - acc: 0.8718 - val\_loss: 0.4021 - val\_acc: 0.9041  
Epoch 55/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4406 - acc: 0.8796 - val\_loss: 0.4024 - val\_acc: 0.8950  
Epoch 56/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4729 - acc: 0.8552 - val\_loss: 0.4209 - val\_acc: 0.8995  
Epoch 57/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4495 - acc: 0.8738 - val\_loss: 0.4057 - val\_acc: 0.8950  
Epoch 58/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4364 - acc: 0.8738 - val\_loss: 0.4103 - val\_acc: 0.9041  
Epoch 59/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4474 - acc: 0.8806 - val\_loss: 0.4201 - val\_acc: 0.8995  
Epoch 60/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4390 - acc: 0.8796 - val\_loss: 0.4006 - val\_acc: 0.9041  
Epoch 61/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4417 - acc: 0.8679 - val\_loss: 0.4103 - val\_acc: 0.8858  
Epoch 62/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4369 - acc: 0.8826 - val\_loss: 0.4490 - val\_acc: 0.8950

Epoch 63/100  
32/32 [=====] - 1s 33ms/step - loss: 0.4370 - acc: 0.8796 - val\_loss: 0.4044 - val\_acc: 0.9132  
Epoch 64/100  
32/32 [=====] - 1s 33ms/step - loss: 0.4373 - acc: 0.8767 - val\_loss: 0.4322 - val\_acc: 0.9087  
Epoch 65/100  
32/32 [=====] - 1s 33ms/step - loss: 0.4548 - acc: 0.8748 - val\_loss: 0.4583 - val\_acc: 0.8813  
Epoch 66/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4433 - acc: 0.8806 - val\_loss: 0.4005 - val\_acc: 0.9041  
Epoch 67/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4320 - acc: 0.8875 - val\_loss: 0.4629 - val\_acc: 0.8767  
Epoch 68/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4681 - acc: 0.8640 - val\_loss: 0.4255 - val\_acc: 0.8950  
Epoch 69/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4410 - acc: 0.8855 - val\_loss: 0.4092 - val\_acc: 0.8904  
Epoch 70/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4491 - acc: 0.8699 - val\_loss: 0.4062 - val\_acc: 0.8858  
Epoch 71/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4365 - acc: 0.8836 - val\_loss: 0.4255 - val\_acc: 0.8950  
Epoch 72/100  
32/32 [=====] - 1s 37ms/step - loss: 0.4479 - acc: 0.8757 - val\_loss: 0.4006 - val\_acc: 0.9041  
Epoch 73/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4367 - acc: 0.8855 - val\_loss: 0.4015 - val\_acc: 0.9041  
Epoch 74/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4501 - acc: 0.8748 - val\_loss: 0.4137 - val\_acc: 0.8904  
Epoch 75/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4315 - acc: 0.8806 - val\_loss: 0.4017 - val\_acc: 0.8904  
Epoch 76/100  
32/32 [=====] - 1s 36ms/step - loss: 0.4370 - acc: 0.8777 - val\_loss: 0.4138 - val\_acc: 0.8858  
Epoch 77/100  
32/32 [=====] - 1s 39ms/step - loss: 0.4377 - acc: 0.8689 - val\_loss: 0.4402 - val\_acc: 0.8904  
Epoch 78/100  
32/32 [=====] - 1s 37ms/step - loss: 0.4523 - acc: 0.8650 - val\_loss: 0.4253 - val\_acc: 0.8904  
Epoch 79/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4266 - acc: 0.8865 - val\_loss: 0.4118 - val\_acc: 0.8904  
Epoch 80/100  
32/32 [=====] - 1s 35ms/step - loss: 0.4554 - acc: 0.8601 - val\_loss: 0.4442 - val\_acc: 0.8813  
Epoch 81/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4386 - acc: 0.8826 - val\_loss: 0.4245 - val\_acc: 0.9041  
Epoch 82/100  
32/32 [=====] - 1s 34ms/step - loss: 0.4420 - acc: 0.8777 - val\_loss: 0.4241 - val\_acc: 0.8904  
Epoch 83/100  
32/32 [=====] - 1s 33ms/step - loss: 0.4267 - acc: 0.8845 - val\_loss:

```

0.4136 - val_acc: 0.8813
Epoch 84/100
32/32 [=====] - 1s 34ms/step - loss: 0.4456 - acc: 0.8796 - val_loss:
0.4179 - val_acc: 0.8995
Epoch 85/100
32/32 [=====] - 1s 36ms/step - loss: 0.4520 - acc: 0.8767 - val_loss:
0.4077 - val_acc: 0.8904
Epoch 86/100
32/32 [=====] - 1s 36ms/step - loss: 0.4342 - acc: 0.8914 - val_loss:
0.4059 - val_acc: 0.8904
Epoch 87/100
32/32 [=====] - 1s 37ms/step - loss: 0.4374 - acc: 0.8826 - val_loss:
0.4054 - val_acc: 0.9132
Epoch 88/100
32/32 [=====] - 1s 35ms/step - loss: 0.4331 - acc: 0.8699 - val_loss:
0.4354 - val_acc: 0.8904
Epoch 89/100
32/32 [=====] - 1s 35ms/step - loss: 0.4439 - acc: 0.8708 - val_loss:
0.4048 - val_acc: 0.8904
Epoch 90/100
32/32 [=====] - 1s 40ms/step - loss: 0.4469 - acc: 0.8620 - val_loss:
0.4204 - val_acc: 0.8995
Epoch 91/100
32/32 [=====] - 1s 35ms/step - loss: 0.4405 - acc: 0.8757 - val_loss:
0.4212 - val_acc: 0.8858
Epoch 92/100
32/32 [=====] - 1s 35ms/step - loss: 0.4440 - acc: 0.8787 - val_loss:
0.4305 - val_acc: 0.9087
Epoch 93/100
32/32 [=====] - 1s 34ms/step - loss: 0.4679 - acc: 0.8640 - val_loss:
0.4439 - val_acc: 0.8858
Epoch 94/100
32/32 [=====] - 1s 34ms/step - loss: 0.4471 - acc: 0.8738 - val_loss:
0.4126 - val_acc: 0.8858
Epoch 95/100
32/32 [=====] - 1s 36ms/step - loss: 0.4377 - acc: 0.8826 - val_loss:
0.4181 - val_acc: 0.8858
Epoch 96/100
32/32 [=====] - 1s 33ms/step - loss: 0.4450 - acc: 0.8777 - val_loss:
0.4032 - val_acc: 0.9178
Epoch 97/100
32/32 [=====] - 1s 35ms/step - loss: 0.4243 - acc: 0.8855 - val_loss:
0.4019 - val_acc: 0.9132
Epoch 98/100
32/32 [=====] - 1s 34ms/step - loss: 0.4309 - acc: 0.8885 - val_loss:
0.4014 - val_acc: 0.9132
Epoch 99/100
32/32 [=====] - 1s 36ms/step - loss: 0.4465 - acc: 0.8855 - val_loss:
0.4026 - val_acc: 0.9041
Epoch 100/100
32/32 [=====] - 1s 35ms/step - loss: 0.4382 - acc: 0.8855 - val_loss:
0.4068 - val_acc: 0.8904

```


We'll now plot the loss and accuracy graphs for Model 3. You'll notice that the loss is a lot higher at the start, and that's because we've changed our loss function. To plot such that the window is zoomed in between 0 and 1.2 for the loss, we add an additional line of code (`plt.ylim`) when plotting

```

In [28]: plt.plot(hist_3.history['loss'])
plt.plot(hist_3.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')


```

```
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.ylim(top=1.2, bottom=0)
plt.show()
```



```
In [29]: plt.plot(hist_3.history['acc'])
plt.plot(hist_3.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='lower right')
plt.show()
```





As compared to Model 2, you should see that there's less overfitting!